

Society as a System in Pursuit of Outcomes: An Essay

Honggao Cao
September 2026

Abstract

A team meets its target and is asked to do more next year. A school rises in a ranking and begins organizing its work around staying there. A good result becomes a reason to repeat a practice, even when questions about it remain unanswered. This essay follows such familiar situations to ask how the pursuit of outcomes comes to organize social life. Outcomes help us recognize achievement and hold people responsible for what their work accomplishes. But the signs we use to recognize success can also narrow attention, redirect effort, raise expectations, and strengthen arrangements that become difficult to reconsider. Hospital care, education, organizational performance, and economic progress show why a better number and a better situation are not always the same thing. The concern is not that society wants results. It is that we can become increasingly effective at pursuing what counts as success while losing sight of what we wanted the success to make possible.

Keywords: outcomes; social expectations; performance; comparison; institutions; progress

1. We Met the Target

Imagine a team that has spent a demanding year trying to complete more customer orders. At the year-end meeting, the report shows that the team has exceeded its target. The manager thanks everyone. There is relief, perhaps a celebration. Then the conversation turns to next year.

Now that the team has shown what it can do, a higher target seems reasonable.

Nobody has to intend anything unkind for this to happen. The manager sees evidence of greater capability. The team wants to do well. Those responsible for the organization want it to improve. Yet the achievement has done two things at once: it has rewarded the past year's effort and increased what will be expected of the next year's effort.

Suppose someone in the room asks what the increase meant for customers. Perhaps their orders arrived sooner, or people who previously went without the service could now use it. That would be something to celebrate. But how was the increase achieved? A better way of working could make it sustainable. Longer hours or postponed maintenance could make it harder to repeat. The final total does not tell us which changes made the improvement possible or what they cost.

The meeting can move on without resolving that difference. The target has been met, and everyone knows what that means for the report. What it means for the work, the people doing it, and the people it serves may remain less clear.

We can recognize versions of this situation well beyond a workplace. A grade once greeted with delight can become the grade a family expects next time. An organization can similarly turn an exceptional year into the minimum it expects to maintain. Achievement brings satisfaction, but it can also change the standard against which the next achievement will be judged.

There are reasons to want these outcomes. A qualification can open a door. A profitable year can keep a business operating. A shorter wait for treatment can make a difficult time less burdensome. The concern begins when reaching the next visible result becomes easier to explain than the reason for wanting it. We can say what the target is, who met it, and how the result compares with last year. We may find it harder to say what would make further pursuit worthwhile.

Consider a parent looking at a report card. A grade gives a quick answer to one question: how did the child perform in the assessed work? The parent may also want to know whether the child understands the subject, has become curious about it, or needs help. The grade is useful, but the conversation can become much narrower if it begins and ends with whether the mark went up.

The same narrowing can occur when a hospital reports its performance or a government announces that the economy has grown. We need to know what was accomplished. We also need to remember what we hoped those accomplishments would do for people. A report can help us ask that question, or it can become a reason to stop asking it.

This is the concern behind looking at society as a system in pursuit of outcomes. The word society can sound distant, but the pursuit takes place in ordinary decisions: what parents praise, what employers reward, what customers compare, what institutions report, and what those responsible for them ask for next. These decisions affect one another. A result matters partly because other people use it, and their use of it gives us reasons to pursue it.

The team in the meeting is therefore not simply choosing how ambitious to be. It is working within expectations that its own success can help raise. To understand that pressure, we need to begin with something less troubling: why looking at outcomes is so useful in the first place.

2. Why We Want an Answer We Can See

Suppose someone close to you needs hospital treatment. You may want to know how long patients usually wait, whether the treatment helps, and how often serious complications occur. You are unlikely to be satisfied with a statement that the hospital is committed to excellent care. You want to know what happens to people who go there.

That request is reasonable. The purpose of the work includes helping patients, and patients have reasons to ask whether it helps. The same applies when parents ask whether children are learning, when a customer asks whether a product works, or when citizens ask whether a public service delivers what was promised. An account of effort and good intentions cannot answer every question about what was achieved.

The difficulty is that much of the work takes place outside our view. A family cannot observe every treatment at competing hospitals. A governing board cannot accompany every patient from diagnosis to recovery. Reported outcomes give them access to information they would otherwise lack.

A waiting-time figure, for example, summarizes how long a group of patients waited. To compare hospitals, we need the figures to cover comparable services and use the same definition of waiting. With that common basis, we can notice a difference and ask why it exists. The comparison becomes possible

because we have selected something the hospitals can report in the same way, leaving much of the detail of each case aside [1].

Leaving detail out is not necessarily a fault. A summary that repeated every case would no longer do the job for which we wanted the summary. The trouble is forgetting which question the summary can answer. Knowing that patients waited less does not yet tell us everything about their care.

A mortality rate makes the difficulty especially clear. A hospital treating unusually ill patients may have a higher rate of deaths than one treating less severe cases. To decide what the difference says about care, we need to understand the patients' conditions and what was done for them. The outcome has to be considered alongside the care and circumstances that preceded it [2].

The distinction is ordinary once we keep the people in view. A patient recovered. A report summarizes recovery among many patients. Someone then concludes that one hospital provides better care. The report helps with that conclusion, but it does not eliminate the additional questions. What happened, how we describe it, and what we conclude from the description are connected steps, not the same step repeated in different words.

A score of 3–1 settles who won a match, but not which team will win next time. A hospital report likewise needs to be read in relation to the particular question we are asking.

Reports also let people responsible for an institution ask questions of those running it. Suppose a hospital's waits keep growing. Its board can ask managers what is causing the delay and what they have tried. The figures do not supply the explanation, but they prevent the conversation from resting entirely on reassurance. Work on auditing and the use of indicators in public decisions describes how reports become part of these arrangements for oversight [3], [4].

From there, a request for improvement follows naturally. The board may ask managers to reduce waiting times and provide resources for the work. A specific target gives people a shared aim and a way to check their progress. Under appropriate conditions, goals of this kind can focus attention, sustain effort, and improve performance [5]. The report is no longer only an account of the past. It helps direct what people will do next.

This is an important change, even though none of its steps seems surprising. We first wanted a number because we could not observe all the work. We then used it to compare institutions, ask for explanations, and decide what should improve. The people doing the work now know that the number matters to their future. They have good reasons to pay attention to it. The next question is what that attention changes.

3. When the Score Starts Choosing the Work

Imagine a student deciding how to spend the evening before an examination. There is a question in the subject that has caught her interest, but she has also been told which material is likely to be tested. Time is short. She studies the tested material.

Nothing about this choice requires dishonesty or a lack of curiosity. The examination matters. It may affect a grade, admission, or an opportunity she wants. Yet the expected assessment has helped decide

what she learns tonight. If similar choices recur, it can help shape her education well before the final result is reported.

The point is not that studying for an examination prevents learning. It can encourage exactly the learning that was intended. The difficulty appears when doing well in the assessment and understanding the subject begin to pull apart. A student can then have a good reason to pursue the grade while leaving something valuable unlearned.

Institutions face a similar choice. A university may want to provide a good education and also want a high ranking. Those aims can support each other, but they are not identical. Once a published rank affects how prospective students and others view the institution, university leaders have reasons to examine what determines it.

A study of law-school rankings found changes in expectations, admissions practices, resource allocation, reporting, and other organizational decisions as schools responded to the rankings [6]. The ranking was intended to describe schools. Schools then changed in response to how they were described. The next ranking reflected institutions already adapting to the previous one.

To see why the distinction matters, suppose a school has money for one of two improvements. One would help students in a way that its ranking recognizes. The other would meet a need that the ranking does not record. Both may be worthwhile. But when the school's reputation is tied to its rank, the first has an additional advantage: its benefit will be easier to show to the people judging the school.

Repeated choices of that kind can change an institution without anyone declaring that education no longer matters. What can be reported begins to receive more attention because it is easier to defend, fund, and reward. The unreported need has not become less important to the student experiencing it. It has become harder to place in the institutional argument about what should be done.

This is more subtle than cheating. A school may improve the dimensions rewarded by the ranking while leaving other educational needs behind. Alternatively, it may find ways to improve its position without an equivalent improvement for students. Those are different responses. A higher rank does not, by itself, identify which has occurred.

The same pressure can affect a worker whose activity is counted. Suppose the worker is expected to complete a certain number of cases. A difficult case needs extra attention, but spending that time means completing fewer cases today. The worker is now balancing the person's need against the reported measure of a good day's work. We should be cautious about interpreting a response to that pressure as a purely personal failure. The way the work is judged has helped create the choice.

The response can also make the measure less useful to those reading it. Once people learn to produce the measured feature specifically, it may tell us less about the broader activity we hoped it would describe. This is the concern behind established warnings about turning measures into targets [7], [8, pp. 34–35]. We need to look at what people do to meet the target, rather than assume that every improvement in the number means the same thing.

There are also cases in which the pressure helps. In the early 2000s, England's National Health Service used waiting-time targets, public reporting, and managerial consequences to address delays. Research comparing England with Scotland found that the English target regime reduced the share of patients

facing long waits for planned hospital treatment [9]. Other research documented gaming risks and the danger of allowing the measured part to stand for the whole of care [10]. Both findings belong in the story.

For a patient whose wait became shorter, the improvement was not an empty achievement. It changed an experience that mattered. We lose that fact if we treat every response to a target as manipulation. We lose a different fact if the shorter wait is allowed to settle every question about the service. A useful measure can promote improvement in one part of the work while leaving other parts to be examined.

The question is therefore concrete: what did people change in order to produce the better figure, and whom did that change help? Looking closely at the work can reveal a benefit, a displacement of attention, or some combination. Merely knowing that the target was met cannot complete the inquiry.

Even when the outcome is unquestionably valuable, however, a further difficulty remains. People can achieve something worthwhile and learn the wrong lesson from having achieved it. Success itself can become reassuring in ways the evidence does not justify.

4. The Reassurance of Success

“It worked last time” can be a sensible reason to try something again. Experience would be of little use if we learned nothing from it. But the phrase can also close a discussion before we have understood what last time actually showed.

Suppose a piece of equipment develops a fault but continues operating. The completed work is good news. The fault is not. If we attend only to the completed work, we may begin to treat the fault as less concerning simply because it has not yet stopped us. Repetition can strengthen that reassurance: we have seen the same problem several times, and the equipment has kept working.

The history preceding the Challenger accident gives this ordinary argument a terrible significance. On earlier shuttle flights, rubber O-rings in the solid rocket boosters had shown erosion and signs of hot gas moving farther into the joints than intended. The seals were meant to prevent that escape. Yet the missions had returned safely. NASA and contractor Morton Thiokol were therefore confronting warning evidence and successful flights at the same time [11, Vol. I, ch. 6].

A safe return established something important: the mission had been completed without losing the shuttle. It did not explain away the damage to the seals. One could conclude that the mission had succeeded despite an unresolved problem. One could instead take the successful flight as reassurance that the observed damage was acceptable. Those readings of the same event pointed toward different choices about the next flight.

The Rogers Commission found that NASA and Thiokol gradually accepted erosion and hot-gas blow-by as an acceptable flight risk [11, Vol. I, ch. 6]. What had first appeared as a departure from intended behavior became part of the accumulated record of flight experience. The record contained repeated success, and that success helped make the departure more familiar and less alarming.

Vaughan's study places this history within the organization's culture, routines, production pressures, and divided responsibilities [12]. This was not simply a matter of one person overlooking a warning. People

worked within established ways of judging evidence. As those judgments were repeated, an explanation that had once required discussion could become part of what later decisions took for granted.

Organizations can carry lessons into their routines even when the experiences behind them leave the reasons for success uncertain [13].

By the eve of the Challenger launch, unusually low temperature had become a major concern. Thiokol engineers initially recommended against launching with O-ring temperatures below 53°F. Management later reversed that recommendation despite continuing engineering objections [11, Vol. I, ch. 5]. The warning was not something that became visible only after the accident. People were raising it before the decision was made.

Challenger launched on January 28, 1986, and was destroyed seventy-three seconds later. The failure of the field joint in the right solid rocket booster was the immediate technical cause [11, Vol. I, ch. 4].

The lesson is not that successful experience tells us nothing. It is that the conclusion must fit the experience. “The mission returned safely” and “the system is safe for the next launch” are different claims. The first cannot do all the work required to establish the second.

We make less dramatic versions of that move when we judge people after learning whether their decisions worked. In experiments, people rated otherwise identical decisions more favorably when the outcomes were favorable, even though the information available to the original decision-maker was unchanged [14]. Knowing the ending affected how the earlier choice was judged.

An institution then acts on that judgment. It rewards a manager, repeats a practice, relaxes a safeguard, or gives an approach more resources. The interpretation of success becomes part of the setting in which the next decision is made. A mistaken lesson can therefore matter long after the first successful event.

Imagine an employee raising a recurring problem during a year of strong performance. A reply that the organization is doing well may sound decisive. But it answers a different question. The employee may be asking about a weakness that has not yet caused a visible failure. Good results should help the organization investigate with some confidence and resources. They should not make the question harder to ask.

This is one danger in allowing success to speak for everything around it. Risks can remain unresolved beneath an accomplishment. Other costs can remain unseen because they fall on people whose experience is absent from the report. To understand a good year, we have to ask whose year it was.

5. A Good Year for Whom?

“The company had a good year.”

Suppose that means its earnings rose. The statement is useful, and those earnings may support jobs, investment, and future activity. Yet different people can have had quite different years within that company. Employees may have gained security, or they may have worked longer under greater pressure. Customers may have received better service, or the company may have reduced services they valued. A financial result cannot tell us which of these stories occurred.

A favorable report need not conceal something bad. It still leaves us to ask how the improvement was achieved and how its benefits and burdens were shared.

The same issue arises when we hear that average income has increased. The average can rise while some people gain a great deal and others gain little or lose income. None of that requires the reported average to be wrong. It means that an average and an account of everyone's circumstances are different kinds of information.

A similar concern motivates efforts to look beyond gross domestic product, or GDP, when discussing social progress. GDP measures economic production. A fuller account asks how material gains are shared, what happens to quality of life, and whether improvements can be sustained [15], [16]. We can recognize an increase in production while still asking what it has changed in people's lives.

Imagine listening to a report that conditions have improved while your own circumstances have become more difficult. It is possible that you are one of the people the average does not describe well. Your experience does not disprove the aggregate figure. The aggregate figure does not disprove your experience either. A serious discussion needs room for both.

Time introduces a similar difficulty. Suppose a company raises this year's earnings by postponing maintenance or cutting an investment whose benefits would arrive later. The current financial gain may be accurately reported. Whether it represents a lasting improvement remains open. Part of the cost may simply have been moved beyond the period being discussed.

Conversely, an organization may spend money now on training, repairs, research, or spare capacity that makes it better prepared for later demands. Those activities can look expensive beside work that produces an immediate return. Their purpose becomes clearer when we follow what they enable or prevent over a longer period.

The report ends on a date. The consequences do not. That difference can matter greatly to people working within a reporting cycle. Someone defending a maintenance budget may have to explain a future problem that has not happened. Someone proposing a spending cut can point to a saving that will appear in the next report. The second argument is easier to display, though it need not be the better decision.

The choices also reach beyond the people present when they are made. A later manager may inherit the neglected equipment. Future customers may encounter the delay. Residents may live with a public service whose capacity was reduced years earlier. Asking whether an improvement lasted is therefore also a way of asking who eventually paid for it.

Even after the facts are clearer, people may disagree about the tradeoff. Some may accept a present burden for a long-term benefit; others may question whether the burden is fairly shared or the benefit likely to arrive. A measure can inform that disagreement. It cannot decide every question about what a society should value.

Adding more measures can help, but a larger report still leaves choices about what to include and how to present it [16]. We need neither one perfect number nor a description so exhaustive that nobody can use it. We need to keep the question attached to the answer: what improved, for whom, at what cost, and for how long?

Those questions should remain available even when the news is good. Otherwise the favorable headline can move directly into the next decision. The organization has succeeded, so its method is repeated, its targets rise, and the people associated with it gain influence. What was left out of the account of success may then be left out of the future it helps create.

6. Why Success Asks for More

Return to the team that exceeded its target. Before the successful year, reaching that level looked ambitious. Afterward, it is part of what everyone knows the team can do. The work may have changed little, but the expectations surrounding it have changed.

Suppose the following year is also good, though not quite as good as the exceptional one. The team might now have to explain a disappointment that would once have been celebrated. The point of comparison has moved. Organizations assess their performance against both their own history and what others have achieved, so an earlier success can change what later looks satisfactory [17].

A past success can therefore become a future requirement. That is one reason achievement does not always bring the sense of arrival people expect. There may be a brief celebration, followed by a revised understanding of what counts as ordinary performance.

Sometimes the higher expectation is justified. Better equipment or greater experience may make an old target too easy. People who still need a service may benefit when an organization expands. But an exceptional year might also have depended on circumstances that will not recur, or on effort that cannot be sustained. Treating the achievement as a permanent increase in capability skips the question of how it was produced.

Other people's achievements can raise the pressure further. Imagine two organizations competing for the same customers. One improves its service; the other has a reason to respond to remain attractive. Its response then gives the first a reason to improve again. This continuing exchange is one way competition can drive organizational learning and change [18]. The work need not stop when either organization reaches its earlier goal, because the situation against which that goal was set has changed.

There can be substantial benefits in that competition. Customers may receive better service, and both organizations may learn. But keeping one's position and making the underlying activity better are still different accomplishments. In a ranking of the same schools, every school could improve its teaching without every school rising in rank. Conversely, one school could rise because others declined, without improving its own teaching at all.

That difference helps explain why the pursuit can become exhausting. Where success means staying ahead, improvement by others creates a fresh demand rather than a shared reason for satisfaction. A qualification that once distinguished an applicant loses that power when everyone has it. Acquiring it may still teach something valuable, but its value as a way to stand out has changed. The individual has a reason to seek a further distinction even if the group as a whole has gained little from the additional competition for position.

Advising each person to be less competitive does not remove the pressure. An applicant cannot decide alone that employers will stop valuing the distinction. A university cannot decide alone that families will stop consulting rankings. In some social arrangements, continuing to grow or do more becomes a condition of maintaining the position already reached [19]. People can keep participating for understandable reasons while collectively sustaining demands that no participant chose in full.

Competition does more than change what participants are expected to achieve. Success within it can also change the resources available for their next attempt. A highly ranked university may attract stronger applicants. A company with a profitable record may obtain funding more readily. A manager associated with a successful project may be trusted with a larger one. The next round begins with opportunities partly shaped by the previous round.

It can be reasonable to trust demonstrated performance. Where success brings such advantages, however, later performance is no longer a fresh test under the same conditions. The successful organization may have gained resources, attention, and freedom that others lack. Repeated success is then partly a record of what those advantages made possible, not simply a series of independent tests of the same original capability.

As people invest in the arrangement, changing it becomes harder. An institution acquires staff and equipment, develops routines, and builds relationships around the approach that brought success. People learn how to work within it and make plans that assume it will continue. Using what they have built can become increasingly attractive compared with switching to an approach that requires fresh investment and coordination [20].

There is value in that stability. Starting again continually would waste learning and make cooperation difficult. Yet the same commitments can keep an institution attached to an approach after its purpose or circumstances have changed. A question about whether the approach still helps may be heard as a threat to the work and investments already built around it.

We can now see how the effort to meet a target can become a continuing demand to do more. Achievement raises expectations. Comparison encourages further effort. Recognition draws resources toward some people and institutions. Those resources support activities and commitments that make the approach harder to leave. Each step has reasons behind it, but their combined effect can keep the pursuit going without a fresh decision about what it is for.

That is why society's attachment to outcomes cannot be understood only as a collection of personal ambitions. We reward one another, depend on one another's judgments, and respond to standards that other people also feel required to meet. A change in direction may require more than an individual deciding to want less. It may require the people setting, using, and responding to the standards to reconsider them together.

7. What Did We Want to Make Better?

A missed target usually prompts a search for better ways to reach it. The harder question—especially after success—is whether that target still represents what the work most needs to accomplish.

Return to the hospital board. Suppose it finds that long waits are caused by poor scheduling. Improving the scheduling addresses the problem while leaving the objective intact. Now suppose the board discovers something different: managers have become so occupied with the reported waiting time that other needs receive too little attention. Asking them simply to improve the same number faster would leave that problem in place.

The board would need to reconsider how it judges the service. Waiting time could remain important without carrying the whole judgment about care. Staff would need to know that reporting an unmeasured difficulty is useful rather than an admission that they have failed to focus on the target. The change would concern what the institution asks people to notice, explain, and improve.

These are different ways to learn from experience: finding a better way to reach the target, or reconsidering the assumptions that made it the target [21], [22]. The practical difference is easy to miss when every disappointing result produces another demand for better execution. Sometimes people need a better way to do the work. Sometimes they need permission to question what has come to count as doing it well.

The distinction matters even after success. A team that has met its target may be in a good position to ask what the achievement has made possible. Should it take on more work? Improve a neglected part of the service? Repair equipment, develop skills, or make the pace more sustainable? The target's completion creates an opportunity for judgment. It need not answer the question of what comes next before that judgment begins.

At the level of society, reconsideration becomes more difficult because people do not share every priority. More economic production may be urgently needed where basic needs remain unmet. Elsewhere, the question may be how its benefits are shared, what damage accompanies it, or whether it improves the security and possibilities of ordinary lives. Recognizing those differences does not supply one agreed definition of progress. It prevents a growth figure from deciding the definition on everyone's behalf.

Looking beyond GDP matters when the additional information changes what receives attention and support [15], [16]. If a report reveals that some people have not benefited from growth, the next decision should address their circumstances rather than simply demand a higher growth figure.

The same modest change is possible in a conversation about a school or a workplace. We can acknowledge the good grade and ask what the child has come to understand. We can welcome improved performance and ask whether customers experienced the benefit. We can take a successful year seriously without using it to dismiss warnings or silence questions about how the work was done.

Doing this requires more than adding a caution at the bottom of a report. When rewards, resources, and reputations depend almost entirely on one number, people receive a clear instruction regardless of what the accompanying statement of purpose says. Keeping a broader purpose alive means allowing it to affect actual decisions: what work is supported, which concerns receive a hearing, and what people are asked to explain.

The team at the opening need not leave its meeting with a lower target. Further improvement may be both possible and valuable. But it should leave with a clearer reason for the next target than the fact that the

last one was exceeded. The people served by the work, the conditions under which it is done, and the consequences that have not yet appeared should have a place in that decision.

A society needs to accomplish things. It also needs to remain able to ask which accomplishments are worth pursuing and what they are doing to the people and institutions pursuing them. More can be better. A better score can record progress. Neither connection should become something we stop examining simply because the pursuit is familiar, successful, or difficult to interrupt.

We can celebrate the result and still ask what it means. We can value achievement without making every achievement a reason to demand still more of the same. The question is not whether the work should produce outcomes. It is whether, while becoming better at producing them, we still remember what we wanted to make better.

AI-use disclosure. Generative AI assisted with literature discovery, drafting, analytical and editorial review, and document preparation across versions of this essay. The author directs the work and its purpose and retains responsibility for the paper's claims, citations, and conclusions.

References

- [1] Espeland, W. N., & Stevens, M. L. (1998). Commensuration as a Social Process. *Annual Review of Sociology*, 24, 313-343. DOI: 10.1146/annurev.soc.24.1.313.
- [2] Donabedian, A. (1988). The Quality of Care: How Can It Be Assessed? *JAMA*, 260(12), 1743-1748. DOI: 10.1001/jama.1988.03410120089033.
- [3] Power, M. (1997). *The Audit Society: Rituals of Verification*. Oxford: Oxford University Press.
- [4] Davis, K. E., Kingsbury, B., & Merry, S. E. (2012). Indicators as a Technology of Global Governance. *Law & Society Review*, 46(1), 71-104. DOI: 10.1111/j.1540-5893.2012.00473.x.
- [5] Locke, E. A., & Latham, G. P. (2002). Building a Practically Useful Theory of Goal Setting and Task Motivation: A 35-Year Odyssey. *American Psychologist*, 57(9), 705-717. DOI: 10.1037/0003-066X.57.9.705.
- [6] Espeland, W. N., & Sauder, M. (2007). Rankings and Reactivity: How Public Measures Recreate Social Worlds. *American Journal of Sociology*, 113(1), 1-40. DOI: 10.1086/517897.
- [7] Goodhart, C. A. E. (2006). The ECB and the Conduct of Monetary Policy: Goodhart's Law and Lessons from the Euro Area. *JCMS: Journal of Common Market Studies*, 44(4), 757-778. DOI: 10.1111/j.1468-5965.2006.00661.x.
- [8] Campbell, D. T. (2011). Assessing the Impact of Planned Social Change. *Journal of MultiDisciplinary Evaluation*, 7(15), 3-43. DOI: 10.56645/jmde.v7i15.297. Authorized reprint of the 1976 paper.
- [9] Propper, C., Sutton, M., Whitnall, C., & Windmeijer, F. (2008). Did 'Targets and Terror' Reduce Waiting Times in England for Hospital Care? *The B.E. Journal of Economic Analysis & Policy*, 8(2), Article 5. DOI: 10.2202/1935-1682.1863.
- [10] Bevan, G., & Hood, C. (2006). What's Measured Is What Matters: Targets and Gaming in the English Public Health Care System. *Public Administration*, 84(3), 517-538. DOI: 10.1111/j.1467-9299.2006.00600.x.
- [11] Presidential Commission on the Space Shuttle Challenger Accident. (1986). Report to the President by the Presidential Commission on the Space Shuttle Challenger Accident. Washington, DC: U.S. Government Printing Office.
- [12] Vaughan, D. (1996). *The Challenger Launch Decision: Risky Technology, Culture, and Deviance at NASA*. Chicago: University of Chicago Press.

- [13] Levitt, B., & March, J. G. (1988). Organizational Learning. *Annual Review of Sociology*, 14, 319-338. DOI: 10.1146/annurev.so.14.080188.001535.
- [14] Baron, J., & Hershey, J. C. (1988). Outcome Bias in Decision Evaluation. *Journal of Personality and Social Psychology*, 54(4), 569-579. DOI: 10.1037/0022-3514.54.4.569.
- [15] Stiglitz, J. E., Sen, A., & Fitoussi, J.-P. (2009). Report by the Commission on the Measurement of Economic Performance and Social Progress. Paris: Commission on the Measurement of Economic Performance and Social Progress.
- [16] Stiglitz, J. E., Fitoussi, J.-P., & Durand, M. (2018). Beyond GDP: Measuring What Counts for Economic and Social Performance. Paris: OECD Publishing. DOI: 10.1787/9789264307292-en.
- [17] Greve, H. R. (1998). Performance, Aspirations, and Risky Organizational Change. *Administrative Science Quarterly*, 43(1), 58-86. DOI: 10.2307/2393591.
- [18] Barnett, W. P., & Hansen, M. T. (1996). The Red Queen in Organizational Evolution. *Strategic Management Journal*, 17(S1), 139-157. DOI: 10.1002/smj.4250171010.
- [19] Rosa, H., Dörre, K., & Lessenich, S. (2017). Appropriation, Activation and Acceleration: The Escalatory Logics of Capitalist Modernity and the Crises of Dynamic Stabilization. *Theory, Culture & Society*, 34(1), 53-73. DOI: 10.1177/0263276416657600.
- [20] Pierson, P. (2000). Increasing Returns, Path Dependence, and the Study of Politics. *American Political Science Review*, 94(2), 251-267. DOI: 10.2307/2586011.
- [21] Argyris, C. (1977). Double Loop Learning in Organizations. *Harvard Business Review*, 55(5), 115-125.
- [22] Moynihan, D. P. (2005). Goal-Based Learning and the Future of Performance Management. *Public Administration Review*, 65(2), 203-216. DOI: 10.1111/j.1540-6210.2005.00445.x.