

Society as a System in Pursuit of Outcomes: An Essay

Honggao Cao
September 2026

Abstract

A team meets its target and is asked to do more next year. A school rises in a ranking and begins organizing its work around staying there. A good result becomes a reason to repeat a practice, even when questions about it remain unanswered. This essay follows such familiar situations to ask how the pursuit of outcomes comes to organize social life. Outcomes help us recognize achievement and hold people responsible for what their work accomplishes. But the signs we use to recognize success can also narrow attention, redirect effort, raise expectations, and strengthen arrangements that become difficult to reconsider. Hospital care, education, organizational performance, and economic progress show why a better number and a better situation are not always the same thing. The concern is that we can become increasingly effective at pursuing what counts as success while losing sight of what we wanted the success to make possible.

Keywords: outcomes; social expectations; performance; comparison; institutions; progress

1. We Met the Target

Imagine a team that has spent a demanding year trying to complete more customer orders. At the year-end meeting, the report shows that the team has exceeded its target. The manager thanks everyone. There is relief, perhaps a celebration. Then the conversation turns to next year.

Now that the team has shown what it can do, a higher target seems reasonable.

The manager reads the result as evidence of greater capability. The team wants to do well, and those responsible for the organization want it to improve. These understandable responses give the achievement two consequences: recognition for the past year's effort and a higher expectation for the next year's effort.

Suppose someone in the room asks what the increase meant for customers. Perhaps their orders arrived sooner, or people who previously went without the service could now use it. That would be something to celebrate. But how was the increase achieved? A better way of working could make it sustainable. Longer hours or postponed maintenance could make it harder to repeat. The final total does not tell us which changes made the improvement possible or what they cost.

The meeting can move on without resolving that difference. Meeting the target gives everyone a clear result to report. Using that result to set the next target requires a further judgment about the work, the people doing it, and the people it serves. If that judgment is skipped, the higher expectation can be established before anyone knows whether the improvement can be sustained.

We can recognize versions of this situation well beyond a workplace. A grade once greeted with delight can become the grade a family expects next time. An organization can similarly turn an exceptional year

into the minimum it expects to maintain. Achievement brings satisfaction, but it can also change the standard against which the next achievement will be judged.

There are reasons to want these outcomes. A qualification can open a door. A profitable year can keep a business operating. A shorter wait for treatment can make a difficult time less burdensome. The concern begins when reaching the next visible result becomes easier to explain than the reason for wanting it. We can say what the target is, who met it, and how the result compares with last year. We may find it harder to say what would make further pursuit worthwhile.

Consider a parent looking at a report card. A grade gives a quick answer to one question: how did the child perform in the assessed work? The parent may also want to know whether the child understands the subject, has become curious about it, or needs help. The grade is useful, but the conversation can become much narrower if it begins and ends with whether the mark went up.

The same narrowing can occur when a hospital reports its performance or a government announces that the economy has grown. We need to know what was accomplished. We also need to remember what we hoped those accomplishments would do for people. A report can help us ask that question, or it can become a reason to stop asking it.

This is the concern behind looking at society as a system in pursuit of outcomes. The word society can sound distant, but the pursuit takes place in ordinary decisions: what parents praise, what employers reward, what customers compare, what institutions report, and what those responsible for them ask for next. These decisions affect one another. A result matters partly because other people use it, and their use of it gives us reasons to pursue it.

The team in the meeting is therefore not simply choosing how ambitious to be. It is working within expectations that its own success can help raise. To understand that pressure, we need to begin with something less troubling: why looking at outcomes is so useful in the first place.

2. Why We Want an Answer We Can See

Suppose someone close to you needs hospital treatment. You may want to know how long patients usually wait, whether the treatment helps, and how often serious complications occur. You are unlikely to be satisfied with a statement that the hospital is committed to excellent care. You want to know what happens to people who go there.

That request is reasonable. The purpose of the work includes helping patients, and patients have reasons to ask whether it helps. The same applies when parents ask whether children are learning, when a customer asks whether a product works, or when citizens ask whether a public service delivers what was promised. An account of effort and good intentions cannot answer every question about what was achieved.

The difficulty is that much of the work takes place outside our view. A family cannot observe every treatment at competing hospitals. A governing board cannot accompany every patient from diagnosis to recovery. Reported outcomes give them access to information they would otherwise lack.

A waiting-time figure, for example, summarizes how long a group of patients waited. To compare hospitals, we need the figures to cover comparable services and use the same definition of waiting. With that common basis, we can notice a difference and ask why it exists. The comparison becomes possible because we have selected something the hospitals can report in the same way, leaving much of the detail of each case aside [1].

Selecting a common measure allows the summary to do its job. Repeating every case would make comparison much harder. To use the summary well, however, we need to keep its question in view: shorter waits tell us something valuable about access to treatment, while other questions about care require other information.

A mortality rate makes the difficulty especially clear. A hospital treating unusually ill patients may have a higher rate of deaths than one treating less severe cases. To decide what the difference says about care, we need to understand the patients' conditions and what was done for them. The outcome has to be considered alongside the care and circumstances that preceded it [2].

The distinction is ordinary once we keep the people in view. A patient recovered. A report summarizes recovery among many patients. Someone then concludes that one hospital provides better care. The report helps with that conclusion, but it does not eliminate the additional questions. What happened, how we describe it, and what we conclude from the description are connected steps, not the same step repeated in different words.

The additional questions become especially important when a report is used to judge those running an institution. Suppose a hospital's waits keep growing. Its board can ask managers what is causing the delay and what they have tried. The figures supply a reason for that inquiry and prevent the conversation from resting entirely on reassurance; the explanation still requires examining the service. Work on auditing and indicators describes how reports enter such arrangements for oversight [3], [4].

Once the board identifies a delay that needs attention, a request for improvement follows naturally. It may ask managers to reduce waiting times and provide resources for the work. A specific target gives people a shared aim and a way to check progress. Under appropriate conditions, goals of this kind can focus attention, sustain effort, and improve performance [5]. The report now has a different role: the figure first used to describe past experience helps define what people are expected to achieve next.

This is an important change, even though none of its steps seems surprising. We first wanted a number because we could not observe all the work. We then used it to compare institutions, ask for explanations, and decide what should improve. The people doing the work now know that the number matters to their future. They have good reasons to pay attention to it. The next question is what that attention changes.

3. When the Score Starts Choosing the Work

Imagine a student deciding how to spend the evening before an examination. There is a question in the subject that has caught her interest, but she has also been told which material is likely to be tested. Time is short. She studies the tested material.

The examination may affect a grade, admission, or an opportunity she wants. Studying the tested material is an understandable response to those stakes. Yet the expected assessment has helped decide what she learns tonight. If similar choices recur, the assessment can help shape her education well before the final result is reported.

Exam preparation can encourage exactly the learning intended. A difficulty arises when doing well in the assessment and understanding the subject begin to pull apart. The student can then have a good reason to pursue the grade while leaving something valuable unlearned. To understand that choice, we need to consider both what she wants to learn and what the assessment rewards.

At the institutional level, the same tension can enter decisions about what receives time and resources. A university may want to provide a good education and also want a high ranking. Those aims can support each other, but they are not identical. Once a published rank affects how prospective students and others view the institution, university leaders have reasons to examine what determines it.

A study of law-school rankings found changes in expectations, admissions practices, resource allocation, reporting, and other organizational decisions as schools responded to the rankings [6]. The ranking was intended to describe schools. Schools then changed in response to how they were described. The next ranking reflected institutions already adapting to the previous one.

To see why the distinction matters, suppose a school has money for one of two improvements. One would help students in a way that its ranking recognizes. The other would meet a need that the ranking does not record. Both may be worthwhile. But when the school's reputation is tied to its rank, the first has an additional advantage: its benefit will be easier to show to the people judging the school.

Repeated choices of that kind can change an institution without anyone declaring that education no longer matters. What can be reported begins to receive more attention because it is easier to defend, fund, and reward. The unreported need has not become less important to the student experiencing it. It has become harder to place in the institutional argument about what should be done.

The choices can produce several different kinds of improvement. A school may improve education in dimensions recognized by the ranking while leaving other needs behind. It may also find ways to improve its position without an equivalent benefit for students. A higher rank records the position; examining the changes behind it tells us which educational benefit, if any, accompanied the rise.

The same pressure can affect a worker whose activity is counted. Suppose the worker is expected to complete a certain number of cases. A difficult case needs extra attention, but spending that time means completing fewer cases today. The worker is now balancing the person's need against the reported measure of a good day's work. We should be cautious about interpreting a response to that pressure as a purely personal failure. The way the work is judged has helped create the choice.

The response can also make the measure less useful to those reading it. Once people learn to produce the measured feature specifically, it may tell us less about the broader activity we hoped it would describe. This is the concern behind established warnings about turning measures into targets [7], [8, pp. 34–35]. We need to look at what people do to meet the target, rather than assume that every improvement in the number means the same thing.

There are also cases in which the pressure helps. In the early 2000s, England's National Health Service used waiting-time targets, public reporting, and managerial consequences to address delays. Research comparing England with Scotland found that the English target regime reduced the share of patients facing long waits for planned hospital treatment [9]. Other research documented gaming risks and the danger of allowing the measured part to stand for the whole of care [10]. Both findings belong in the story.

For a patient whose wait became shorter, the improvement changed an experience that mattered. That benefit gives us a reason to retain waiting time as an important concern. The evidence about gaming gives us a reason to examine how the target was met and what happened elsewhere in the service. Together, the findings show how a useful measure can promote improvement in one part of the work while leaving other parts to be assessed.

The question is therefore concrete: what did people change in order to produce the better figure, and whom did that change help? Looking closely at the work can reveal a benefit, a displacement of attention, or some combination. Merely knowing that the target was met cannot complete the inquiry.

Even when the outcome is unquestionably valuable, however, a further difficulty remains. People can achieve something worthwhile and learn the wrong lesson from having achieved it. Success itself can become reassuring in ways the evidence does not justify.

4. The Reassurance of Success

“It worked last time” can be a sensible reason to try something again. Experience would be of little use if we learned nothing from it. But the phrase can also close a discussion before we have understood what last time actually showed.

Suppose a piece of equipment develops a fault but continues operating. The completed work is good news. The fault is not. If we attend only to the completed work, we may begin to treat the fault as less concerning simply because it has not yet stopped us. Repetition can strengthen that reassurance: we have seen the same problem several times, and the equipment has kept working.

The history preceding the Challenger accident gives this ordinary argument a terrible significance. On earlier shuttle flights, rubber O-ring seals in the solid rocket booster joints had shown erosion and signs of hot gas passing the primary seals. The seals were meant to prevent that passage. Yet the missions had returned safely. NASA and contractor Morton Thiokol were therefore confronting warning evidence and successful flights at the same time [11, Vol. I, ch. 6].

A safe return established something important: the mission had been completed without losing the shuttle. It did not explain away the damage to the seals. One could conclude that the mission had succeeded despite an unresolved problem. One could instead take the successful flight as reassurance that the observed damage was acceptable. Those readings of the same event pointed toward different choices about the next flight.

The Rogers Commission found that NASA and Thiokol gradually accepted erosion and hot-gas blow-by as an acceptable flight risk [11, Vol. I, ch. 6]. What had first appeared as a departure from intended

behavior became part of the accumulated record of flight experience. The record contained repeated success, and that success helped make the departure more familiar and less alarming.

Vaughan's study helps explain how these judgments became established within the organization's culture, routines, production pressures, and divided responsibilities [12]. People assessed the warning evidence through ways of working that had developed over successive flights. As those assessments were repeated, an explanation that had once required discussion could become part of what later decisions took for granted. Research on organizational learning examines how experience enters routines even when the reasons for successful outcomes remain uncertain [13].

The next launch introduced a condition that made the unresolved seal problem especially consequential: unusually low temperature. Cold reduced the rubber seals' ability to respond quickly as the joints moved, raising concern about whether they would seal in time. On the eve of the Challenger launch, Thiokol engineers initially recommended against launching with O-ring temperatures below 53°F. Management later reversed that recommendation despite continuing engineering objections [11, Vol. I, ch. 5]. The concern was being raised before the launch decision, rather than discovered only after the accident.

Challenger launched on January 28, 1986, and was destroyed seventy-three seconds later. The failure of the field joint in the right solid rocket booster was the immediate technical cause [11, Vol. I, ch. 4].

Successful experience is useful when the conclusion fits what the experience establishes. "The mission returned safely" records a completed flight. "The system is safe for the next launch" requires a judgment about the next flight's conditions and the unresolved warning evidence. Keeping those questions separate allows past success to inform the decision without substituting for the assessment it needs.

The temptation to let an outcome settle a broader judgment also appears in less dramatic settings. In experiments, people rated otherwise identical decisions more favorably when the outcomes were favorable, even though the information available to the original decision-maker was unchanged [14]. Knowing the ending affected how the earlier choice was judged. This finding concerns evaluation after an outcome; it offers another reason to distinguish the result from the quality of the decision that preceded it.

When an institution acts on such a judgment, the interpretation has consequences of its own. It may reward a manager, repeat a practice, relax a safeguard, or give an approach more resources. The account of success enters the setting in which the next decision is made, so an inference that reaches beyond the evidence can matter long after the first successful event.

Imagine an employee raising a recurring problem during a year of strong performance. A reply that the organization is doing well may sound decisive. But it answers a different question. The employee may be asking about a weakness that has not yet caused a visible failure. Good results should help the organization investigate with some confidence and resources. They should not make the question harder to ask.

The employee's question shows how a favorable result can leave an important part of the work unresolved. A similar omission occurs when the result summarizes an organization but leaves out differences among the people involved. To understand a good year, we have to ask whose year it was.

5. A Good Year for Whom?

“The company had a good year.” Suppose that means its earnings rose. Those earnings may support jobs, investment, and future activity, but the people within and around the company can have had quite different years. Employees may have gained security, or they may have worked longer under greater pressure. Customers may have received better service, or the company may have reduced services they valued. To understand what the financial improvement meant, we need to ask how it was achieved and how its benefits and burdens were shared.

Moving from a company to an economy enlarges the same question. Average income can rise while some people gain a great deal and others gain little or lose income. The average may accurately describe the overall change while leaving its distribution unclear. Understanding people's circumstances requires that additional information.

A similar concern motivates efforts to look beyond gross domestic product, or GDP, when discussing social progress. GDP measures economic production. A fuller account asks how material gains are shared, what happens to quality of life, and whether improvements can be sustained [15], [16]. We can recognize an increase in production while still asking what it has changed in people's lives.

Imagine listening to a report that conditions have improved while your own circumstances have become more difficult. The difference may arise because the report describes an average that combines your experience with other people's gains. Understanding how both accounts can be accurate gives us a reason to examine the distribution, rather than dismiss either the reported figure or the person's experience.

Following who benefits also takes us beyond the people and costs visible today. Suppose a company raises this year's earnings by postponing routine maintenance or reducing current spending on training. The saving can improve the current report while leaving equipment or staff less prepared for later demands. Whether the gain represents a lasting improvement depends partly on those later consequences.

An organization that maintains the equipment and trains staff pays costs now for benefits that may appear later. The same is true of some spending on research or spare capacity. Those activities can look expensive beside work that produces an immediate return. Their purpose becomes clearer when we follow what they enable or prevent over a longer period.

The report ends on a date. The consequences do not. That difference can matter greatly to people working within a reporting cycle. Someone defending a maintenance budget may have to explain a future problem that has not happened. Someone proposing a spending cut can point to a saving that will appear in the next report. The second argument is easier to display, though it need not be the better decision.

The choices also reach beyond the people present when they are made. A later manager may inherit the neglected equipment. Future customers may encounter the delay. Residents may live with a public service whose capacity was reduced years earlier. Asking whether an improvement lasted is therefore also a way of asking who eventually paid for it.

Knowing who benefits, who bears the costs, and when they appear gives people a better basis for discussing the tradeoff. It does not guarantee agreement. Some may accept a present burden for a long-

term benefit; others may question whether the burden is fairly shared or the benefit likely to arrive. A measure can inform those judgments while leaving the priorities open for discussion.

Adding more measures can help, but a larger report still leaves choices about what to include and how to present it [16]. We need neither one perfect number nor a description so exhaustive that nobody can use it. We need to keep the question attached to the answer: what improved, for whom, at what cost, and for how long?

Keeping those questions open affects more than the interpretation of a completed year. A favorable headline can become the basis for the next decision: the method is repeated, the target rises, and people associated with the success gain influence. If the account omitted the costs or circumstances supporting the achievement, the next decision may proceed without examining them. The team at the opening shows how easily that move can occur.

6. Why Success Asks for More

Return to the team that exceeded its target. Before the successful year, reaching that level looked ambitious. Afterward, it is part of what everyone knows the team can do. The work may have changed little, but the expectations surrounding it have changed.

Suppose the following year is also good, though not quite as good as the exceptional one. The team might now have to explain a disappointment that would once have been celebrated. Its point of comparison has moved. Organizations assess performance against their own history and what others have achieved, so an earlier success can change what later looks satisfactory [17]. The team has reached its old goal, but that achievement has helped establish a new standard for ordinary performance.

Sometimes the higher expectation is justified. Better equipment or greater experience may make an old target too easy. People who still need a service may benefit when an organization expands. But an exceptional year might also have depended on circumstances that will not recur, or on effort that cannot be sustained. Treating the achievement as a permanent increase in capability skips the question of how it was produced.

The team's own history is only one source of expectations. Other organizations' achievements can change what customers expect and what managers regard as attainable. Imagine two organizations competing for the same customers. One improves its service; the other responds to remain attractive. Its response gives the first a reason to improve again. This continuing exchange is one way competition can drive organizational learning and change [18]. Reaching an earlier goal does not end the work when the comparison against which the goal was set keeps changing.

Customers may receive better service as both organizations learn. To judge that benefit, however, we need to distinguish improvement in the service from improvement in relative position. In a ranking of the same schools, every school could improve its teaching without every school rising in rank. Conversely, one school could rise because others declined, without improving its own teaching. The position and the underlying activity can move differently.

That difference helps explain why the pursuit can become exhausting. Where success means staying ahead, improvement by others creates a fresh demand rather than a shared reason for satisfaction. A qualification that once distinguished an applicant loses that power when everyone has it. Acquiring it may still teach something valuable, but its value as a way to stand out has changed. The individual has a reason to seek a further distinction even if the group as a whole has gained little from the additional competition for position.

Advising each person to be less competitive does not remove the pressure. An applicant cannot decide alone that employers will stop valuing the distinction. A university cannot decide alone that families will stop consulting rankings. In some social arrangements, continuing to grow or do more becomes a condition of maintaining the position already reached [19]. People can keep participating for understandable reasons while collectively sustaining demands that no participant chose in full.

A change in relative position can also change the resources available for the next attempt. A highly ranked university may attract stronger applicants. A company with a profitable record may obtain funding more readily. A manager associated with a successful project may be trusted with a larger one. In such cases, the next round begins with opportunities partly shaped by recognition of the previous result.

It can be reasonable to trust demonstrated performance. Where success brings such advantages, however, later performance is no longer a fresh test under the same conditions. The successful organization may have gained resources, attention, and freedom that others lack. Repeated success is then partly a record of what those advantages made possible, not simply a series of independent tests of the same original capability.

The additional resources can support further work, and that work creates commitments that outlast a single target. An institution acquires staff and equipment, develops routines, and builds relationships around the approach associated with its success. People learn to work within the arrangement and make plans that assume it will continue. Using what they have built can become increasingly attractive compared with an alternative requiring fresh investment and coordination [20].

There is value in that stability. Starting again continually would waste learning and make cooperation difficult. Yet the same commitments can keep an institution attached to an approach after its purpose or circumstances have changed. A question about whether the approach still helps may be heard as a threat to the work and investments already built around it.

The opening team's higher target can thus become part of a longer process. Past achievement changes expectations; other people's performance changes the comparison. When recognition also brings resources, those resources support further work and commitments that can make a change of direction harder. Where these links operate, each new decision can be understandable locally while the continuing pursuit receives less scrutiny than the next result.

This helps explain why society's attachment to outcomes reaches beyond individual ambition. We reward one another, depend on one another's judgments, and respond to standards that others also feel required to meet. Reconsidering the pursuit may therefore require the people setting, using, and responding to the

standards to reconsider them together. The question becomes how an institution can preserve that opportunity while continuing to accomplish useful work.

7. What Did We Want to Make Better?

A missed target usually prompts a search for better ways to reach it. The preceding examples give us a reason to ask an additional question, including after success: does the target still represent what the work most needs to accomplish? To act on that question, we need to distinguish improving the method from reconsidering the aim.

Return to the hospital board. Suppose it finds that long waits are caused by poor scheduling. Improving the scheduling addresses the problem while leaving the objective intact. Now suppose the board discovers something different: managers have become so occupied with the reported waiting time that other needs receive too little attention. Asking them simply to improve the same number faster would leave that problem in place.

The board would need to reconsider how it judges the service. Waiting time could remain important without carrying the whole judgment about care. Staff would need to know that reporting an unmeasured difficulty is useful rather than an admission that they have failed to focus on the target. The change would concern what the institution asks people to notice, explain, and improve.

These are different ways to learn from experience: finding a better way to reach the target, or reconsidering the assumptions that made it the target [21], [22]. The practical difference is easy to miss when every disappointing result produces another demand for better execution. Sometimes people need a better way to do the work. Sometimes they need permission to question what has come to count as doing it well.

After success, the same distinction can change what follows. A team that has met its target may be able to take on more work, improve a neglected part of the service, repair equipment, develop skills, or make the pace more sustainable. Choosing among these possibilities requires returning to the purpose of the work. Completing the target creates an opportunity for that judgment; it need not settle the next target in advance.

Extending that judgment from an institution to society brings differences in people's circumstances and priorities into view. More economic production may be urgently needed where basic needs remain unmet. Elsewhere, the question may concern how its benefits are shared, what damage accompanies it, or whether it improves the security and possibilities of ordinary lives. Keeping those needs visible gives us a basis for discussing progress that a growth figure alone cannot supply.

This is why looking beyond GDP matters when the additional information changes what receives attention and support [15], [16]. If a report reveals that some people have not benefited from growth, understanding their circumstances should help guide the next decision. A higher growth figure may be part of the answer, but the report gives us a reason to examine how its benefits would reach them.

The same modest change is possible in a conversation about a school or a workplace. We can acknowledge the good grade and ask what the child has come to understand. We can welcome improved

performance and ask whether customers experienced the benefit. We can take a successful year seriously without using it to dismiss warnings or silence questions about how the work was done.

Doing this requires more than adding a caution at the bottom of a report. When rewards, resources, and reputations depend almost entirely on one number, people receive a clear instruction regardless of what the accompanying statement of purpose says. Keeping a broader purpose alive means allowing it to affect actual decisions: what work is supported, which concerns receive a hearing, and what people are asked to explain.

The team at the opening need not leave its meeting with a lower target. Further improvement may be both possible and valuable. But it should leave with a clearer reason for the next target than the fact that the last one was exceeded. The people served by the work, the conditions under which it is done, and the consequences that have not yet appeared should have a place in that decision.

A society needs to accomplish things. It also needs to remain able to ask which accomplishments are worth pursuing and what they are doing to the people and institutions pursuing them. More can be better. A better score can record progress. Neither connection should become something we stop examining simply because the pursuit is familiar, successful, or difficult to interrupt.

We can celebrate the result and still ask what it means. We can value achievement without making every achievement a reason to demand still more of the same. The question is not whether the work should produce outcomes. It is whether, while becoming better at producing them, we still remember what we wanted to make better.

AI-Use Disclosure

Generative AI assisted with literature discovery, drafting, analytical and editorial review, and document preparation across versions of this essay. The author directs the work and its purpose and retains responsibility for the paper's claims, citations, and conclusions.

References

- [1] Espeland, W. N., & Stevens, M. L. (1998). Commensuration as a Social Process. *Annual Review of Sociology*, 24, 313-343. DOI: 10.1146/annurev.soc.24.1.313.
- [2] Donabedian, A. (1988). The Quality of Care: How Can It Be Assessed? *JAMA*, 260(12), 1743-1748. DOI: 10.1001/jama.1988.03410120089033.
- [3] Power, M. (1997). *The Audit Society: Rituals of Verification*. Oxford: Oxford University Press.
- [4] Davis, K. E., Kingsbury, B., & Merry, S. E. (2012). Indicators as a Technology of Global Governance. *Law & Society Review*, 46(1), 71-104. DOI: 10.1111/j.1540-5893.2012.00473.x.
- [5] Locke, E. A., & Latham, G. P. (2002). Building a Practically Useful Theory of Goal Setting and Task Motivation: A 35-Year Odyssey. *American Psychologist*, 57(9), 705-717. DOI: 10.1037/0003-066X.57.9.705.
- [6] Espeland, W. N., & Sauder, M. (2007). Rankings and Reactivity: How Public Measures Recreate Social Worlds. *American Journal of Sociology*, 113(1), 1-40. DOI: 10.1086/517897.

- [7] Goodhart, C. A. E. (2006). The ECB and the Conduct of Monetary Policy: Goodhart's Law and Lessons from the Euro Area. *JCMS: Journal of Common Market Studies*, 44(4), 757-778. DOI: 10.1111/j.1468-5965.2006.00661.x.
- [8] Campbell, D. T. (2011). Assessing the Impact of Planned Social Change. *Journal of MultiDisciplinary Evaluation*, 7(15), 3-43. DOI: 10.56645/jmde.v7i15.297. Authorized reprint of the 1976 paper.
- [9] Propper, C., Sutton, M., Whitnall, C., & Windmeijer, F. (2008). Did 'Targets and Terror' Reduce Waiting Times in England for Hospital Care? *The B.E. Journal of Economic Analysis & Policy*, 8(2), Article 5. DOI: 10.2202/1935-1682.1863.
- [10] Bevan, G., & Hood, C. (2006). What's Measured Is What Matters: Targets and Gaming in the English Public Health Care System. *Public Administration*, 84(3), 517-538. DOI: 10.1111/j.1467-9299.2006.00600.x.
- [11] Presidential Commission on the Space Shuttle Challenger Accident. (1986). Report to the President by the Presidential Commission on the Space Shuttle Challenger Accident. Washington, DC: U.S. Government Printing Office. <https://www.nasa.gov/history/rogersrep/genindex.htm>
- [12] Vaughan, D. (1996). *The Challenger Launch Decision: Risky Technology, Culture, and Deviance at NASA*. Chicago: University of Chicago Press.
- [13] Levitt, B., & March, J. G. (1988). Organizational Learning. *Annual Review of Sociology*, 14, 319-338. DOI: 10.1146/annurev.so.14.080188.001535.
- [14] Baron, J., & Hershey, J. C. (1988). Outcome Bias in Decision Evaluation. *Journal of Personality and Social Psychology*, 54(4), 569-579. DOI: 10.1037/0022-3514.54.4.569.
- [15] Stiglitz, J. E., Sen, A., & Fitoussi, J.-P. (2009). Report by the Commission on the Measurement of Economic Performance and Social Progress. Paris: Commission on the Measurement of Economic Performance and Social Progress.
- [16] Stiglitz, J. E., Fitoussi, J.-P., & Durand, M. (2018). *Beyond GDP: Measuring What Counts for Economic and Social Performance*. Paris: OECD Publishing. DOI: 10.1787/9789264307292-en.
- [17] Greve, H. R. (1998). Performance, Aspirations, and Risky Organizational Change. *Administrative Science Quarterly*, 43(1), 58-86. DOI: 10.2307/2393591.
- [18] Barnett, W. P., & Hansen, M. T. (1996). The Red Queen in Organizational Evolution. *Strategic Management Journal*, 17(S1), 139-157. DOI: 10.1002/smj.4250171010.
- [19] Rosa, H., Dörre, K., & Lessenich, S. (2017). Appropriation, Activation and Acceleration: The Escalatory Logics of Capitalist Modernity and the Crises of Dynamic Stabilization. *Theory, Culture & Society*, 34(1), 53-73. DOI: 10.1177/0263276416657600.
- [20] Pierson, P. (2000). Increasing Returns, Path Dependence, and the Study of Politics. *American Political Science Review*, 94(2), 251-267. DOI: 10.2307/2586011.
- [21] Argyris, C. (1977). Double Loop Learning in Organizations. *Harvard Business Review*, 55(5), 115-125.
- [22] Moynihan, D. P. (2005). Goal-Based Learning and the Future of Performance Management. *Public Administration Review*, 65(2), 203-216. DOI: 10.1111/j.1540-6210.2005.00445.x.